

# RELATIONAL DATABASE & SQL



# WHY DO WE NEED DATABASES?

Store data. Persistently.

# EXCEL

Let's start with something we all know

# EXAMPLE

Let's store **cities** and their **inhabitants** using Excel. How would you do it?

cities-inhabitants.xlsx

Rechercher dans la feuille

Accueil Mise en page Tableaux Graphiques SmartArt Formules

F1

	A	B	C	D	E	F	G
1		first_name	last_name	age	city		
2		Julian	Dancy	12	London		
3		Pierre	Dupont	48	Paris		
4		Marie	Durand	35	Paris		
5		Victoria	Davis	17	London		
6		Audrey	Lapierre	24	Paris		
7		Angelique	Lefevre	34	Paris		
8		Melissa	Devlin	41	New York		
9							
10							
11							
12							
13							
14							

cities inhabitants +

Mode Normal Prêt

cities-inhabitants.xlsx

Rechercher dans la feuille

Accueil Mise en page Tableaux Graphiques SmartArt Formules

D1

	A	B	C	D	E	F	G
1		name	surface				
2		Paris	105				
3		London	1572				
4		New York	1214				
5							
6							
7							
8							
9							
10							
11							
12							
13							
14							

cities +

Mode Normal Prêt



cities-inhabitants.xlsx

Rechercher dans la feuille

Accueil Mise en page Tableaux Graphiques SmartArt Formules

D1

	A	B	C	D	E	F	G
1	id	name	surface				
2	1	Paris	105				
3	2	London	1572				
4	3	New York	1214				
5							
6							
7							
8							
9							
10							
11							
12							
13							
14							

cities inhabitants +

Mode Normal Prêt

cities-inhabitants.xlsx

Rechercher dans la feuille

Accueil Mise en page Tableaux Graphiques SmartArt Formules

A9

	A	B	C	D	E	F	G
1	id	first_name	last_name	age	city_id		
2	1	Julian	Dancy	12	2		
3	2	Pierre	Dupont	48	1		
4	3	Marie	Durand	35	1		
5	4	Victoria	Davis	17	2		
6	5	Audrey	Lapierre	24	1		
7	6	Angelique	Lefevre	34	1		
8	7	Melissa	Devlin	41	3		
9							
10							
11							
12							
13							
14							

cities inhabitants +

Mode Normal Prêt

# 1:N RELATION (ONE TO MANY)

An inhabitant **belongs to** one city (or has one city)

# EXCEL++

Let's go further

wagon.xlsx

Rechercher dans la feuille

Accueil Mise en page Tableaux Graphiques SmartArt Formules

A1

	A	B	C	D	E	F
1		first_name	last_name	social_security_number	age	
2		George	Abitbol	1 12 34 89 124 123	42	
3		Michel	Hazavanicus	1 94 91 12 123 492	25	
4						
5						
6						
7						
8						
9						
10						
11						
12						
13						
14						

patients +

Mode Normal Prêt

wagon.xlsx

Rechercher dans la feuille

Accueil Mise en page Tableaux Graphiques SmartArt Formules

B5

	A	B	C	D	E	F
1		first_name	last_name	specialty		
2		Sigmund	Freud	Psychology		
3		Henri	Castafolte	Robotics		
4		John	Doe	Cardiag Surgery		
5						
6						
7						
8						
9						
10						
11						
12						
13						
14						

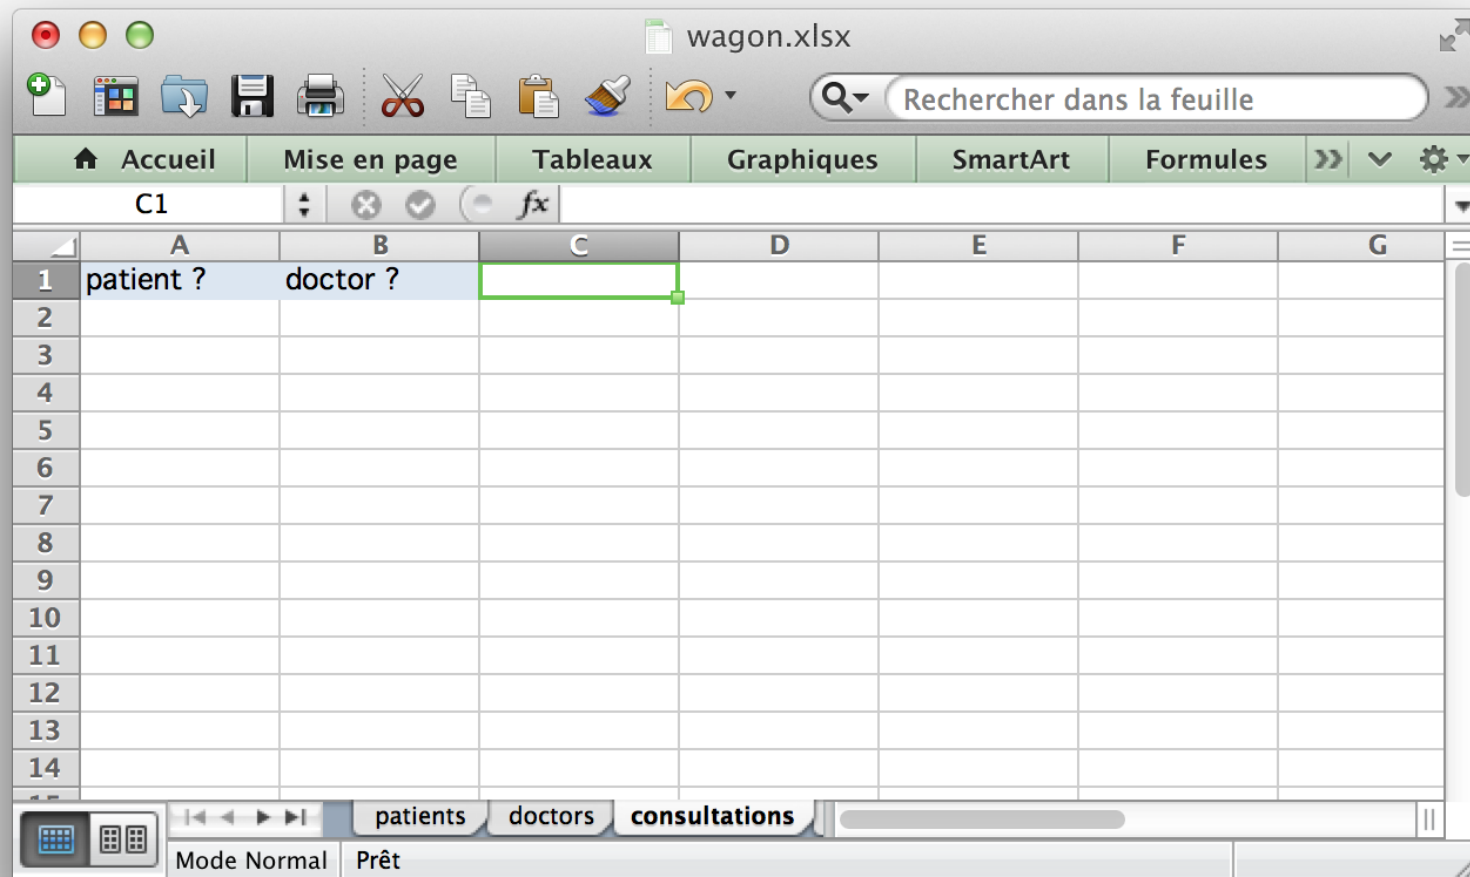
patients doctors +

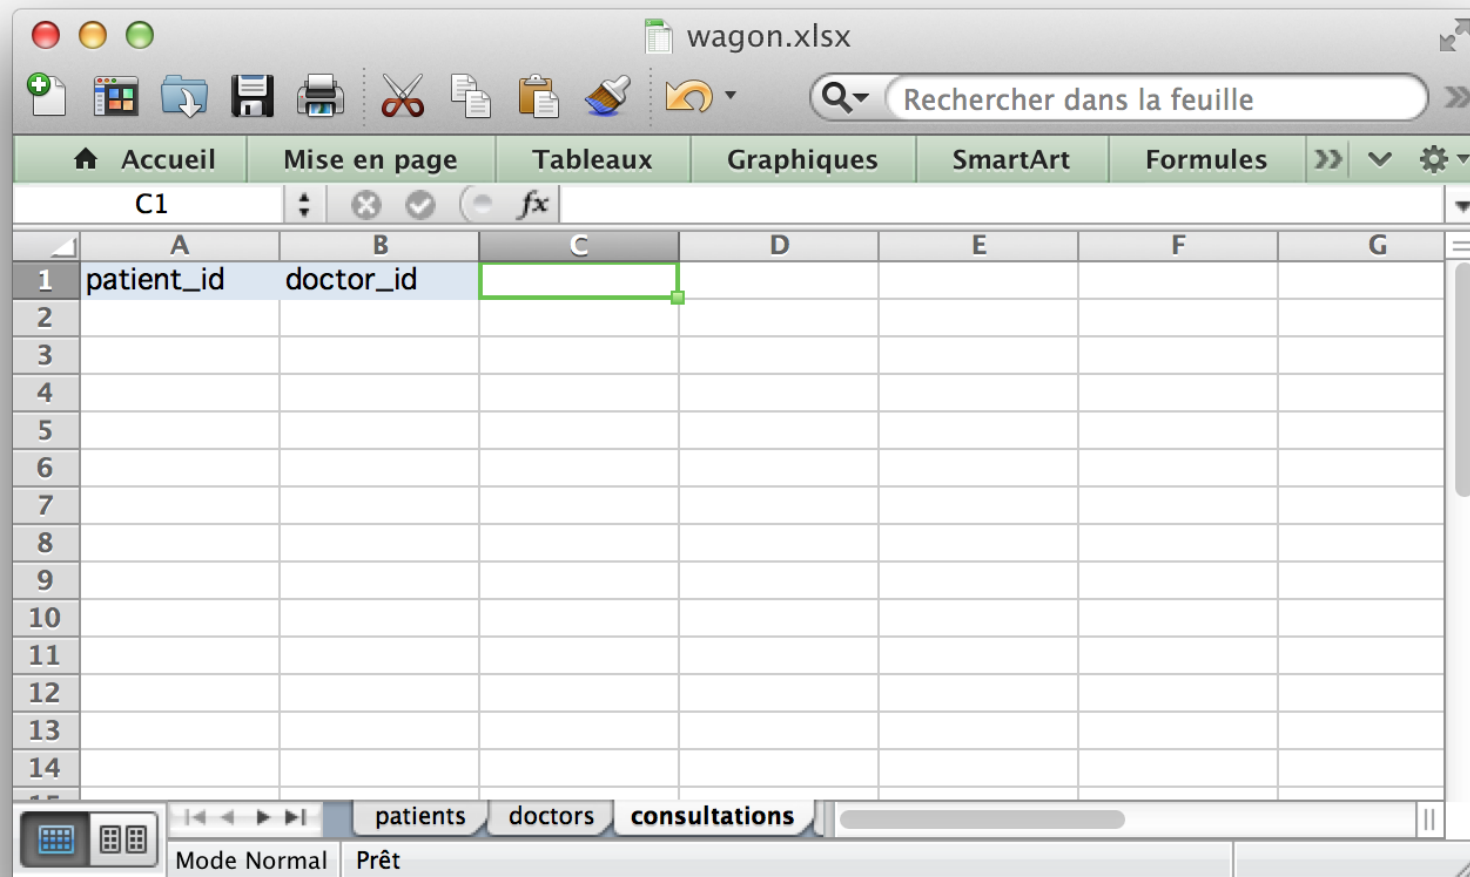
Mode Normal Prêt

**CONSULTATIONS ?**

- **One** doctor can have **many** patients
- **One** patient can see **many** doctors







wagon.xlsx

Rechercher dans la feuille

Accueil Mise en page Tableaux Graphiques SmartArt Formules

A1 id

	A	B	C	D	E	F
1	id	first_name	last_name	social_security_number	age	
2		George	Abitbol	1 12 34 89 124 123	42	
3		Michel	Hazavanicus	1 94 91 12 123 492	25	
4						
5						
6						
7						
8						
9						
10						
11						
12						
13						
14						

patients doctors consultations

Mode Normal Prêt

wagon.xlsx

wagon.xlsx

Rechercher dans la feuille


Accueil Mise en page Tableaux Graphiques SmartArt Formules

A5

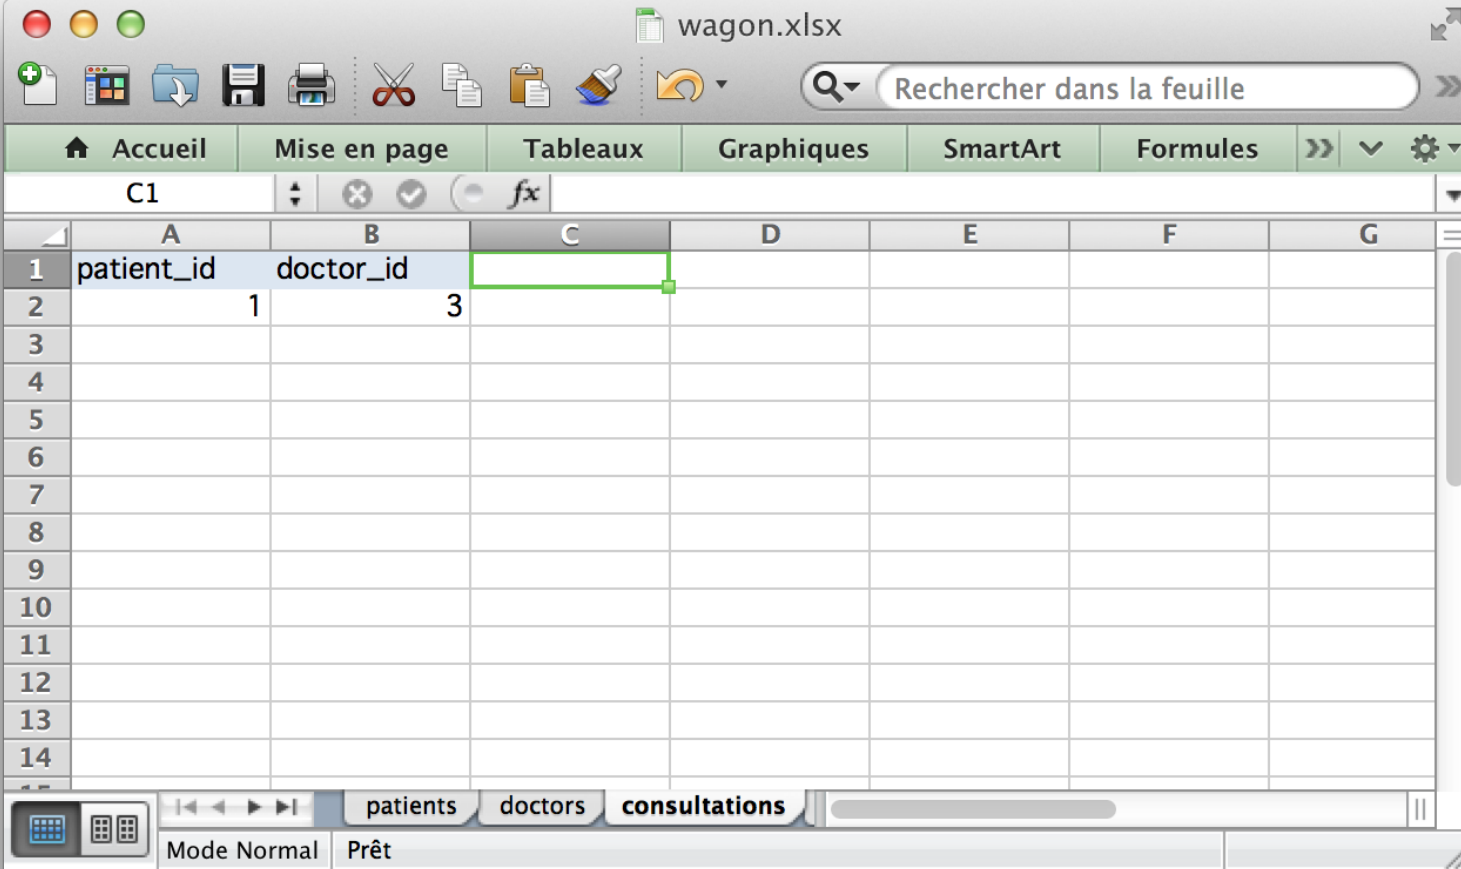
	A	B	C	D	E	F
1	id	first_name	last_name	specialty		
2	1	Sigmund	Freud	Psychology		
3	2	Henri	Castafolte	Robotics		
4	3	John	Doe	Cardiag Surgery		
5						
6						
7						
8						
9						
10						
11						
12						
13						
14						

patients **doctors** consultations

Mode Normal Prêt



George Abitbol (**id** = 1) has seen Doctor John Doe (**id** = 3)



The screenshot shows a Microsoft Excel spreadsheet titled "wagon.xlsx". The spreadsheet has a single sheet named "consultations". The table structure is as follows:

	A	B	C	D	E	F	G
1	patient_id	doctor_id					
2	1	3					
3							
4							
5							
6							
7							
8							
9							
10							
11							
12							
13							
14							

The spreadsheet interface includes a ribbon with tabs: Accueil, Mise en page, Tableaux, Graphiques, SmartArt, and Formules. The status bar at the bottom shows "Mode Normal" and "Prêt".

# N:N RELATION (MANY TO MANY)

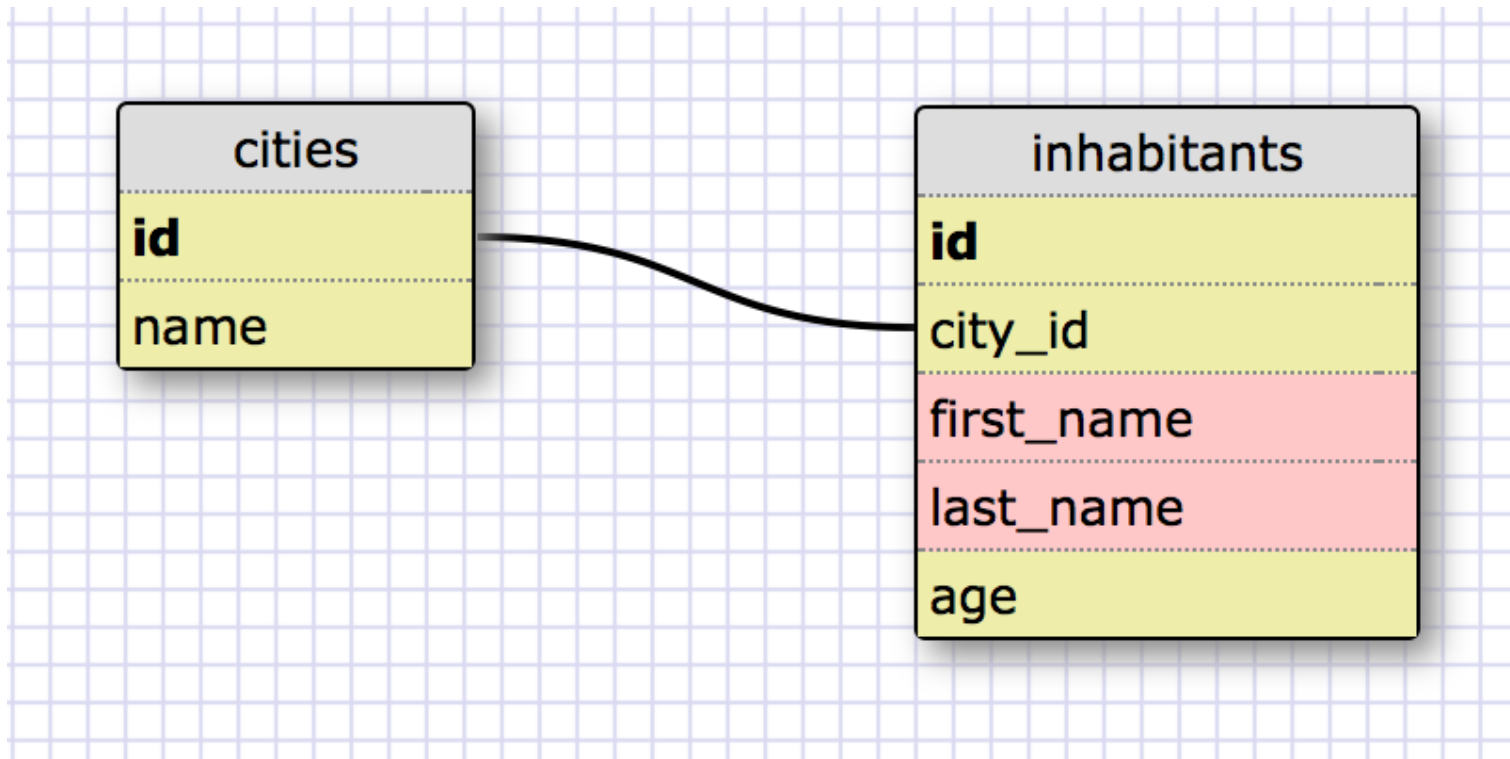
A patient **has many** doctors and a doctor **has many** patients.



You can download this example: [consultations.xlsx](#)

# RELATIONAL DATABASE

1:N

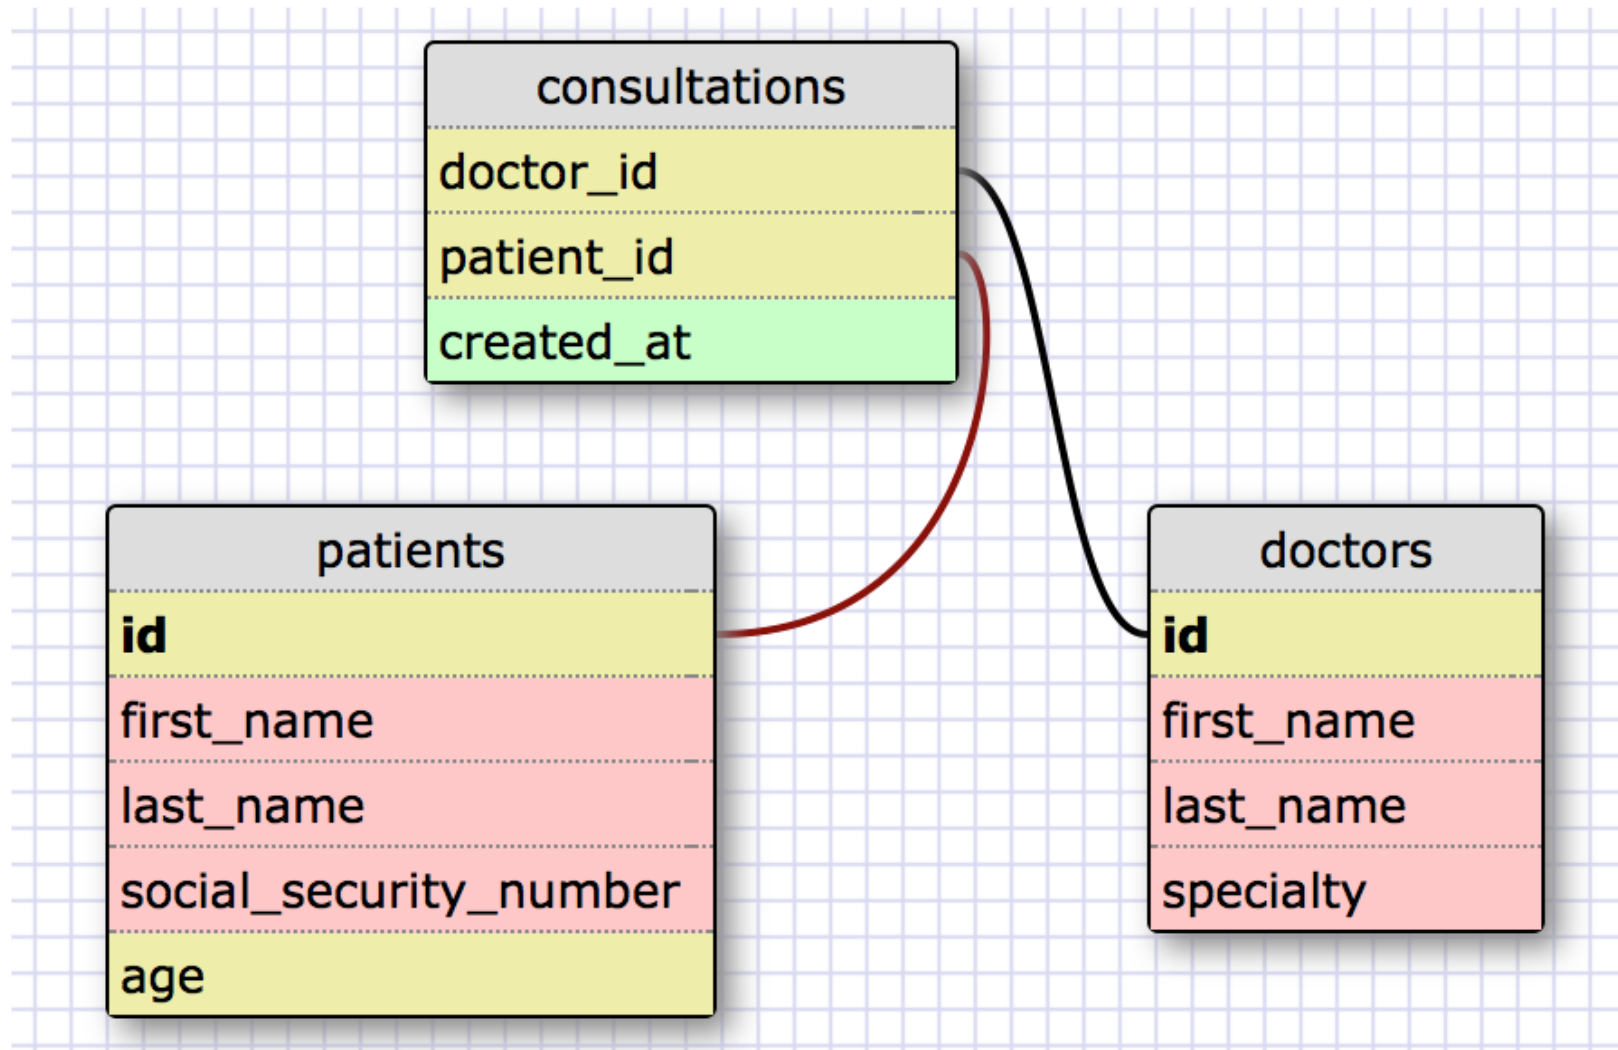


N:N ?

patients
<b>id</b>
first_name
last_name
social_security_number
age

doctors
<b>id</b>
first_name
last_name
specialty

N:N



# VOCABULARY

- A schema is composed of **tables**.
- Each table has a set of **columns**.
- When inserting data in a table, you create a **record** in a new **row**.

# DB SCHEMA COMPOSER

[db.lewagon.com](http://db.lewagon.com)

You can save/load schemas. Try with [patients-doctors.xml](#)



# QUERYING

Retrieve data using the schema.

# SQL

Structured Query Language

# GIVE ME ALL PATIENT NAMES

```
SELECT first_name, last_name FROM patients
```

# GIVE ME ALL DOCTOR NAMES

```
SELECT first_name, last_name FROM doctors
```

# GIVE ME ALL YOU GOT ABOUT PATIENTS

```
SELECT * FROM patients
```

# GIVE ME ALL PATIENTS OF AGE 21

```
SELECT * FROM patients WHERE age = 21
```

# GIVE ME ALL DOCTORS OF CARDIAC SURGERY SPECIALTY

```
SELECT * FROM doctors WHERE specialty = 'Cardiac Surgery'
```

# GIVE ME ALL SURGERY DOCTORS

```
SELECT * FROM doctors WHERE specialty LIKE '%Surgery'
```



# GIVE ME ALL CARDIAC SURGERY DOCTORS NAMED STEVE

```
SELECT * FROM doctors  
WHERE specialty = 'Cardiac Surgery'  
AND first_name = 'Steve'
```

# GIVE ME ALL PATIENTS ORDERED BY AGE

```
SELECT * FROM patients ORDER BY age ASC
```

```
SELECT * FROM patients ORDER BY age DESC
```

# HOW MANY DOCTORS DO I HAVE?

```
SELECT COUNT(*) FROM doctors
```

# COUNT CARDIAC SURGERY DOCTORS

```
SELECT COUNT(*) FROM doctors WHERE specialty = 'Cardiac Surgeon'
```

# COUNT ALL DOCTORS PER SPECIALTY

```
SELECT COUNT(*), specialty  
FROM doctors  
GROUP BY specialty
```

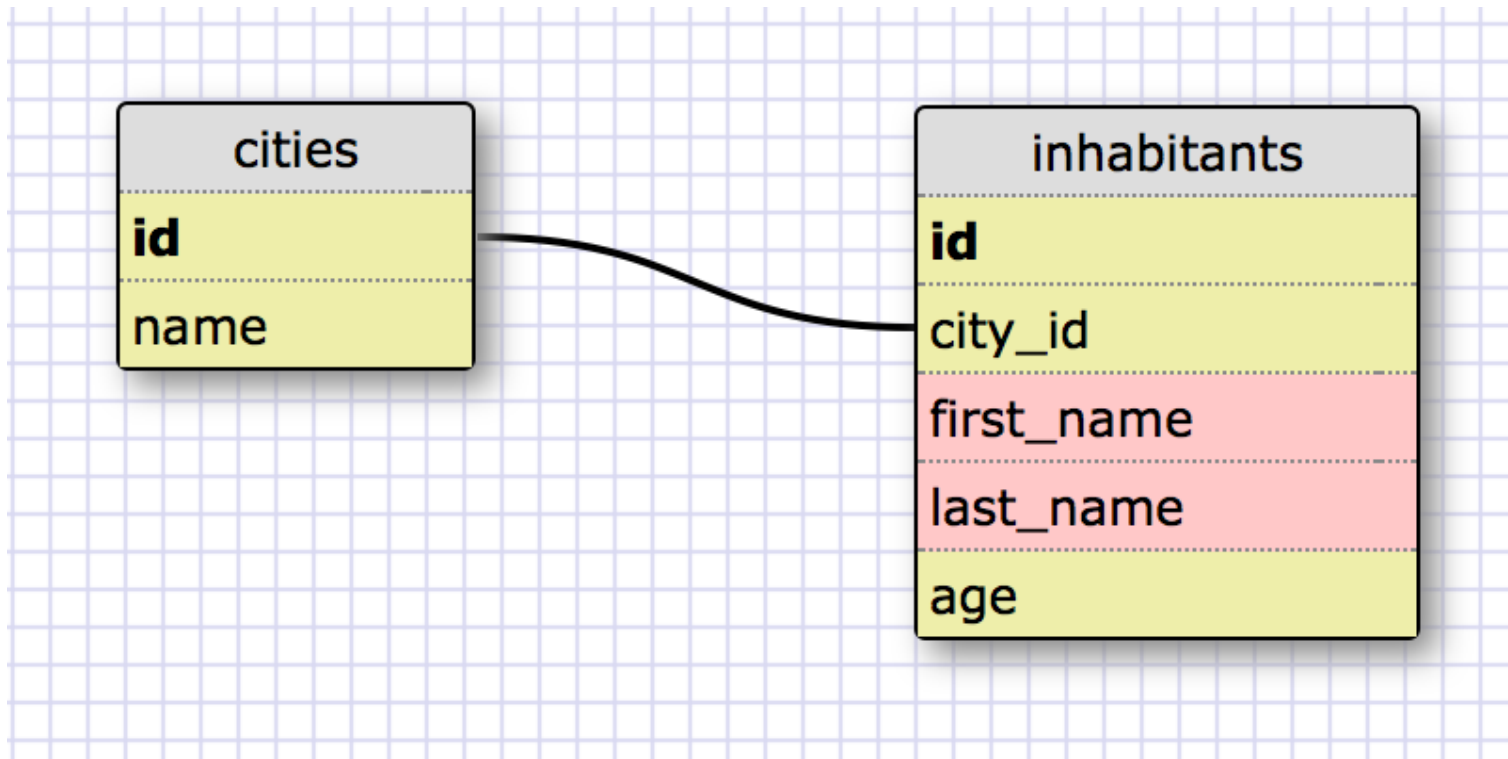
# COUNT ALL DOCTORS PER SPECIALTY, ORDER BY SPECIALTY

You need to rename result column, with AS.

```
SELECT COUNT(*) AS c, specialty  
FROM doctors  
GROUP BY specialty  
ORDER BY c DESC
```

# USING 2 OR MORE TABLES AT ONCE

# GIVEN THIS CITIES/INHABITANTS SCHEMA...





## ... AND THIS DATA

cities		inhabitants					
id	name		id	city_id	first_name	last_name	age
1	Paris		1	3	Sophia	Smith	21
2	Brussels		2	2	Jackson	Williams	30
3	Lille		3	5	Emma	Johnson	29
4	Beirut		4	1	Aiden	Brown	18
5	Bordeaux		5	2	Olivia	Jones	14
			6	4	Liam	Miller	39
			7	4	Ava	Davis	41
			8	1	Lucas	Garcia	43
			9	2	Isabella	Rodriguez	20
			10	1	Noah	Wilson	20

# QUESTION: GIVE ME ALL THE INHABITANTS FROM PARIS

```
SELECT * FROM inhabitants
JOIN cities ON cities.id = inhabitants.city_id
WHERE cities.name = 'Paris'
```

cities			inhabitants				
id	name		id	city_id	first_name	last_name	age
1	Paris		1	3	Sophia	Smith	21
2	Brussels		2	2	Jackson	Williams	30
3	Lille		3	5	Emma	Johnson	29
4	Beirut		4	1	Aiden	Brown	18
5	Bordeaux		5	2	Olivia	Jones	14
			6	4	Liam	Miller	39
			7	4	Ava	Davis	41
			8	1	Lucas	Garcia	43
			9	2	Isabella	Rodriguez	20
			10	1	Noah	Wilson	20
Query Result		<i>"i" is short for "inhabitants"</i>					
cities.id	cities.name	i.id	i.city_id	i.first_name	i.last_name	i.age	
1	Paris	4	1	Aiden	Brown	18	
1	Paris	8	1	Lucas	Garcia	43	
1	Paris	10	1	Noah	Wilson	20	

## QUESTION: GIVE ME ALL THE ADULTS LIVING IN PARIS

```
SELECT * FROM inhabitants  
JOIN cities ON cities.id = inhabitants.city_id  
WHERE inhabitants.age >= 18  
AND cities.name = 'Paris'
```

**GIVEN THIS CONSULTATIONS SCHEMA**

**QUESTION: FOR EACH CONSULTATION, GIVE ME ITS DATE, PATIENT AND DOCTOR NAMES**

```
SELECT c.created_at, p.first_name, p.last_name, d.first_name,  
FROM consultations c  
JOIN patients p ON c.patient_id = p.id  
JOIN doctors d ON c.doctor_id = d.id;
```

# GOING FURTHER

You can read more about INNER (the default one), LEFT OUTER, RIGHT OUTER or FULL OUTER JOIN [here](#), [here](#) and [there](#).

# SQLITE

It is a simple database storing everything in **one** file. Great to quickly test, but not suited for production.



# INSTALLATION

On OSX, run this:

```
brew install sqlite
```

On Ubuntu, run this:

```
sudo apt-get install sqlite3 libsqlite3-dev
```

# QUICK START

Create a new folder, and go into it.  
Create a DB and start typing SQL queries:

```
sqlite3 db.sqlite
```

It will create the `db.sqlite` file.

# TABLE CREATION

```
CREATE TABLE `cities` (  
  `id` INTEGER PRIMARY KEY AUTOINCREMENT,  
  `name` VARCHAR  
);
```

# INSERTING AND QUERYING

```
sqlite> INSERT INTO cities ('name') VALUES ('Paris');  
sqlite> INSERT INTO cities ('name') VALUES ('London');
```

```
sqlite> .mode column  
sqlite> .headers on
```

```
sqlite> SELECT * FROM cities;  
id      name  
-----  
1       Paris  
2       London
```

# HELP

```
sqlite> .help
```

<code>.exit</code>	Exit this program
<code>.header(s) ON OFF</code>	Turn display of headers on or off
<code>.read FILENAME</code>	Execute SQL in FILENAME
<code>.schema ?TABLE?</code>	Show the CREATE statements If TABLE specified, only show tables LIKE pattern TABLE.
<code>.show</code>	Show the current values for various settings
<code>.tables ?TABLE?</code>	List names of tables If TABLE specified, only list tables LIKE pattern TABLE.
<code>[...]</code>	

**IN 2 WEEKS: POSTGRESQL**